

RNNs for Time Series Forecast: Energy Demand Forecast

PART I

1) INTRODUCTION:

The type of deep learning model used in this project is Recurrent Neural Network.

The Concept of Recurrent Neural Networks

Deep learning models called recurrent neural networks (RNNs) are frequently employed to address issues involving sequential input data, such as time series. RNNs are a sort of neural network that can learn from earlier iterations during training because it keeps track of the information it has already processed.

Time Series Forecasting:

A set of techniques in statistics and data science called time series forecasting are used to forecast some variables that vary and evolve over time. Temporal series forecasting's main goal is to predict how target variables will change in the future by looking at past data from a temporal perspective, identifying patterns, and making short- or long-term predictions about change while taking into account the patterns identified.

The models I developed in this project contains LSTM layer.

LSTMs, or Long Short-Term Memory Networks:

When learning a mapping function from inputs to outputs, recurrent neural networks, such as the Long Short-Term Memory network or LSTM, add the explicit handling of order between observations, which is not provided by MLPs or CNNs. These neural networks add native support for input information made up of collections of observations.

In this project we create 3 datasets that is first 50% of data is split for training and intermediate 25% for validation and last 25% for test dataset.

And we will address 2 prediction challenges in this project namely:

3 –hour horizon and 6- hour horizon.

For each of the challenges we will develop:

- a) A 1-input predictor (which will USE VALUES OF JUST *eload* TO PREDICT *eload*).
- b) A 2-input predictor (which will USE VALUES OF BOTH *eload* and *tempf* TO PREDICT *eload*).

2) a) What is the highest (max) value of eload in the TEST SET?

Max eload in the test set: 5224.0

b) What is the lowest (min) value of eload in the TEST SET?

Min eload in the test set: 1979.0

c) what is the difference between the max and the min that you found in parts a) and b) ? (We will call this the “Full Range” of eload, within the TEST SET).

Full range of eload in the test set: 3245.0

Part II-A: 1-Input Predictor (3 hour):

a) sequence_length = 12

b) COMPILE METHOD & FIT METHOD:

```
model.compile(loss='mae' , optimizer=opt , metrics=['mae'])

# Define callbacks using the best model only
best_model_callback = ModelCheckpoint("best_model.h5", monitor='val_loss',
    save_best_only=True)
early_stop = EarlyStopping(monitor='val_loss', patience=5)


# Train model
history = model.fit(train_dataset, epochs=60, batch_size = 128, validation
_data=val_dataset, callbacks=[best_model_callback])
```

c) model.summary()


Model: "sequential"

Layer (type)	Output Shape	Param #
=====	=====	=====
lstm (LSTM)	(None, 64)	16896
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 64)	4160
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
=====	=====	=====
Total params: 23,169		
Trainable params: 23,169		
Non-trainable params: 0		

d) TRAINING AND VALIDATION LOSS



e) TRAINING AND VALIDATION MAE



f) Final Training and Validation LOSS & MAE:

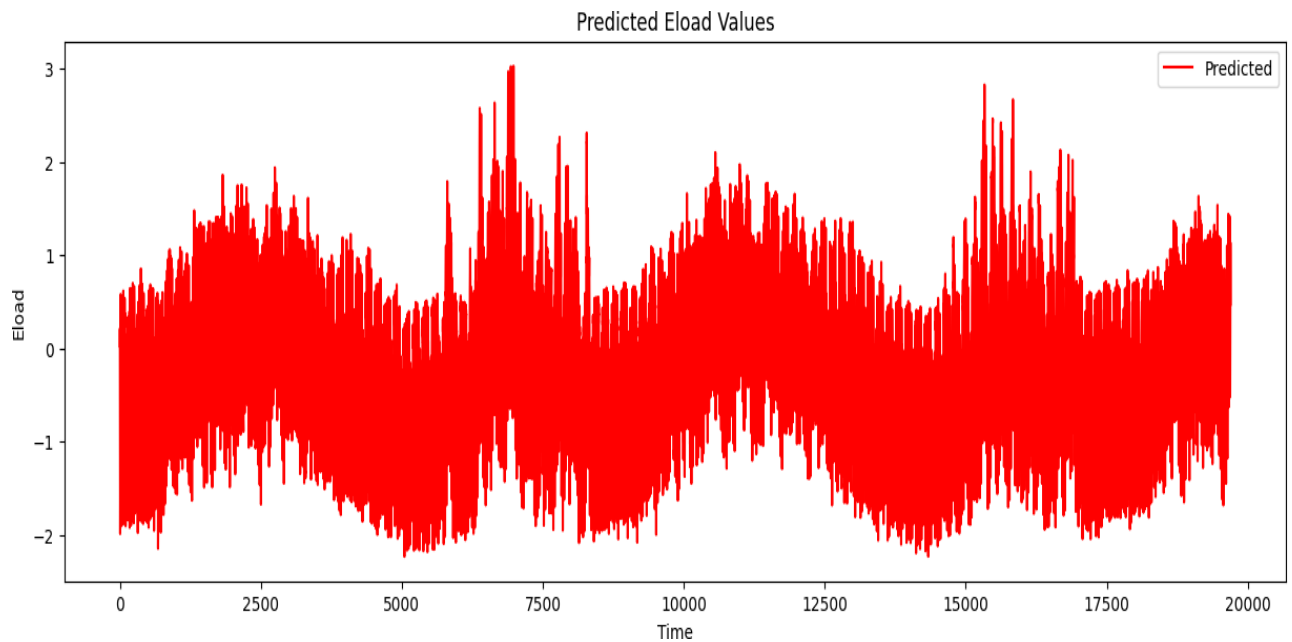
```
loss: 0.2934 - mae: 0.2934 - val_loss: 0.2335 - val_mae: 0.2335
```

```
Training Loss   : 0.2934  
Training MAE    : 0.2934  
Validation LOSS : 0.2335  
Validation MAE  : 0.2335
```


g) MAE obtained on the TEST SET:

```
Test MAE: 0.237017422914505
```


h) Time series plot of all values predicted by the model on the test set:



i) Time series plot of the corresponding targets:



j) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL.

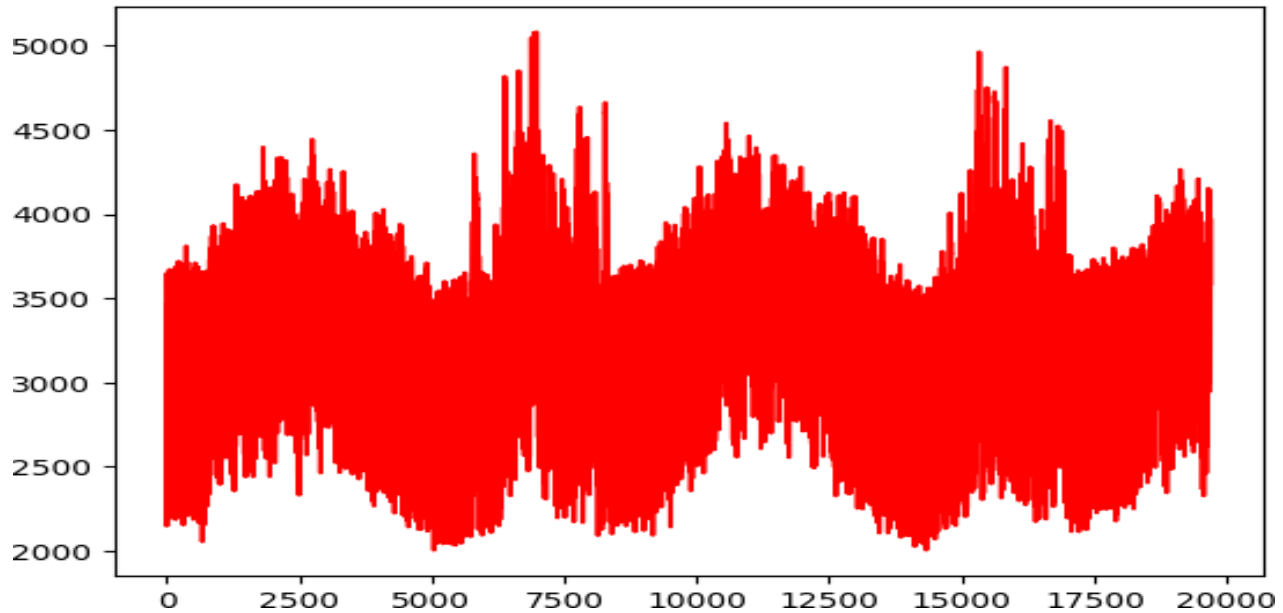


AFTER DOING DENORMALIZATION


g2) True MAE obtained on the TEST SET:

EFFECTIVE real-scale MAE: 198.14

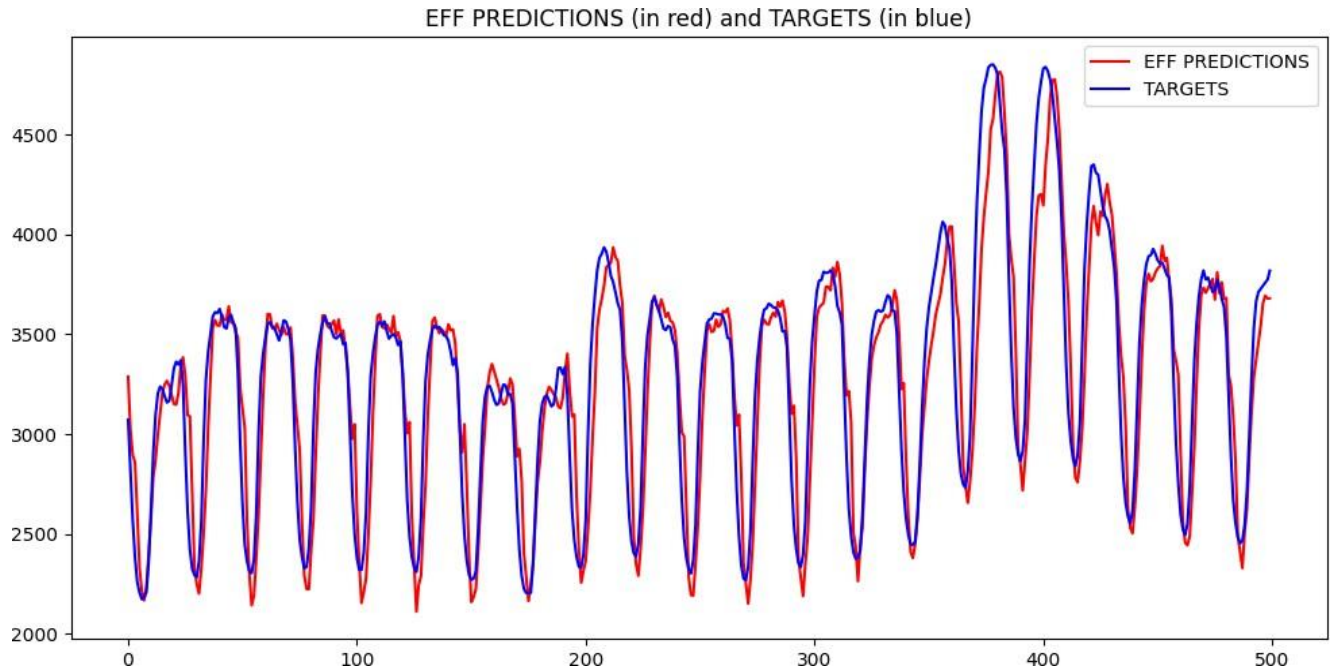
h2) Time series plot of all values predicted by the model on the test set:



i2) Time series plot of the corresponding targets:



j2) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL.



k)

$$PMAE = \frac{(true)MAE \text{ obtained by the model on the test set}}{\text{Full range of elaad in the test set}}$$

$$PMAE = 198.14/3245$$

$$= 0.061$$

$$= 6.1\%$$

Part II-B: 1-Input Predictor (6 hour):

a) sequence_length = 24

b) COMPILE METHOD & FIT METHOD:

```
model.compile(loss='mae' , optimizer=opt , metrics=['mae'])

# Define callbacks using the best model only
best_model_callback = ModelCheckpoint("best_model.h5", monitor='val_loss',
    save_best_only=True)
early_stop = EarlyStopping(monitor='val_loss', patience=5)


# Train model
history = model.fit(train_dataset, epochs=60, batch_size = 128, validation
    _data=val_dataset, callbacks=[best_model_callback])
```

c) model.summary()

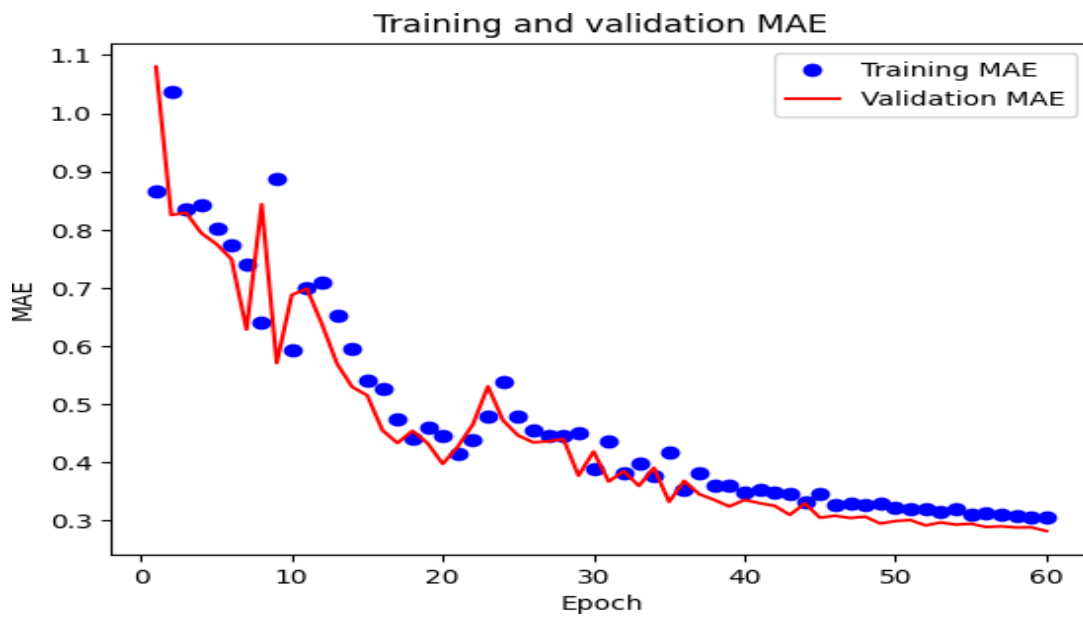
Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 64)	16896
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 64)	4160
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
Total params: 23,169		
Trainable params: 23,169		
Non-trainable params: 0		

d) TRAINING AND VALIDATION LOSS



e) TRAINING AND VALIDATION MAE



f) Final Training and Validation LOSS & MAE:


```
loss: 0.3063 - mae: 0.3063 - val_loss: 0.2820 - val_mae: 0.2820
```

```
Training Loss   : 0.3063  
Training MAE    : 0.3063  
Validation LOSS : 0.2820  
Validation MAE  : 0.2820
```


g) MAE obtained on the TEST SET:

```
Test MAE: 0.283030241727829
```


h) Time series plot of all values predicted by the model on the test set:



i) Time series plot of the corresponding targets



j) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL.




AFTER DOING DENORMALIZATION


g2) True MAE obtained on the TEST SET:

EFFECTIVE real-scale MAE: 192.54

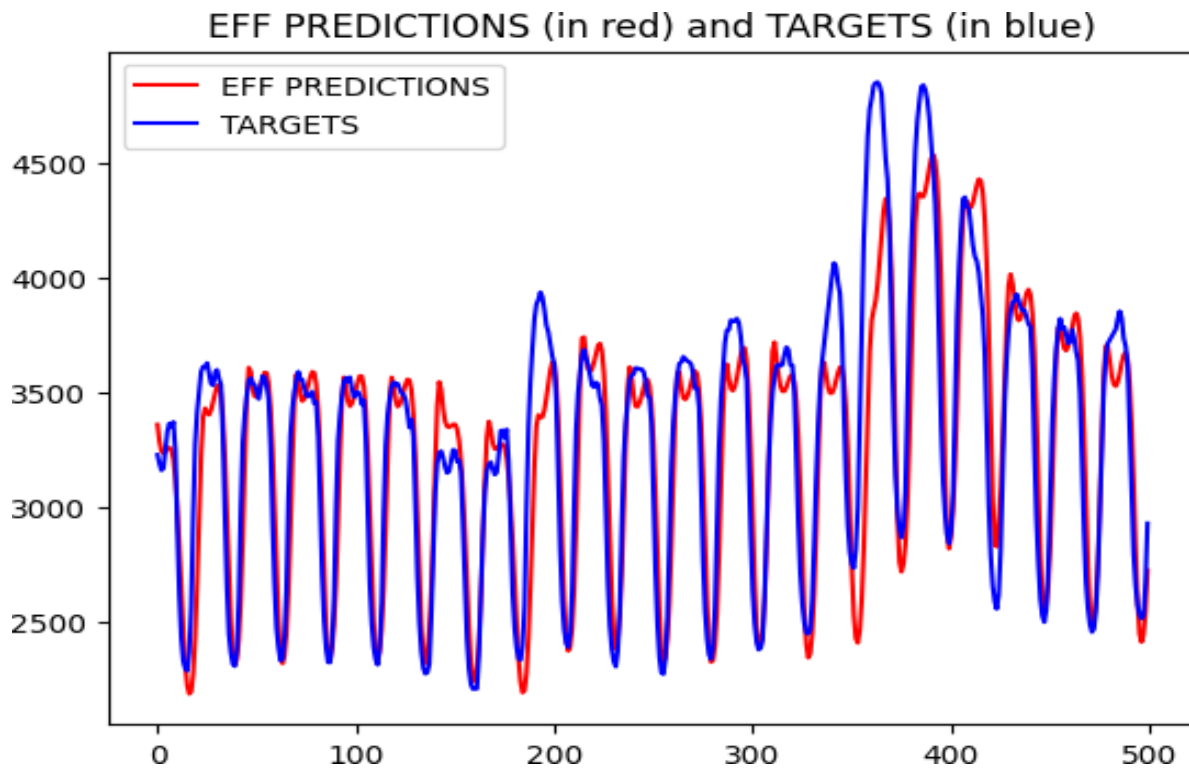
h2) Time series plot of all values predicted by the model on the test set:



i2) Time series plot of the corresponding targets:



j2) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL.



k)

$$PMAE = \frac{(true)MAE \text{ obtained by the model on the test set}}{Full \text{ range of elaad in the test set}}$$

$$PMAE = 192.54/3245$$

$$= 0.059$$

$$= 5.9\%$$

Part III-A: 2-Input Predictor (3 hour):

a) sequence_length = 12

b) COMPILE METHOD & FIT METHOD:

```
model.compile(loss='mae' , optimizer=opt , metrics=['mae'])

# Define callbacks using the best model only
best_model_callback = ModelCheckpoint("best_model.h5", monitor='val_loss',
    save_best_only=True)
early_stop = EarlyStopping(monitor='val_loss', patience=5)


# Train model
history = model.fit(train_dataset, epochs=60, batch_size = 128, validation
_data=val_dataset, callbacks=[best_model_callback])
```

c) model.summary()


Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 64)	17152
dropout_2 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 64)	4160
dropout_3 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 1)	33
Total params: 23,425		
Trainable params: 23,425		
Non-trainable params: 0		

d) TRAINING AND VALIDATION LOSS



e) TRAINING AND VALIDATION MAE



f) Final Training and Validation LOSS & MAE:


```
loss: 0.2616 - mae: 0.2616 - val_loss: 0.2165 - val_mae: 0.2165
```

```
Training Loss   : 0.2616  
Training MAE    : 0.2616  
Validation LOSS  : 0.2165  
Validation MAE   : 0.2165
```


g) MAE obtained on the TEST SET:

```
Test MAE: 0.224911168217659
```


h) Time series plot of all values predicted by the model on the test set:



i) Time series plot of the corresponding targets



j) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL.




AFTER DOING DENORMALIZATION


g2) True MAE obtained on the TEST SET:

EFFECTIVE real-scale MAE: 172.75

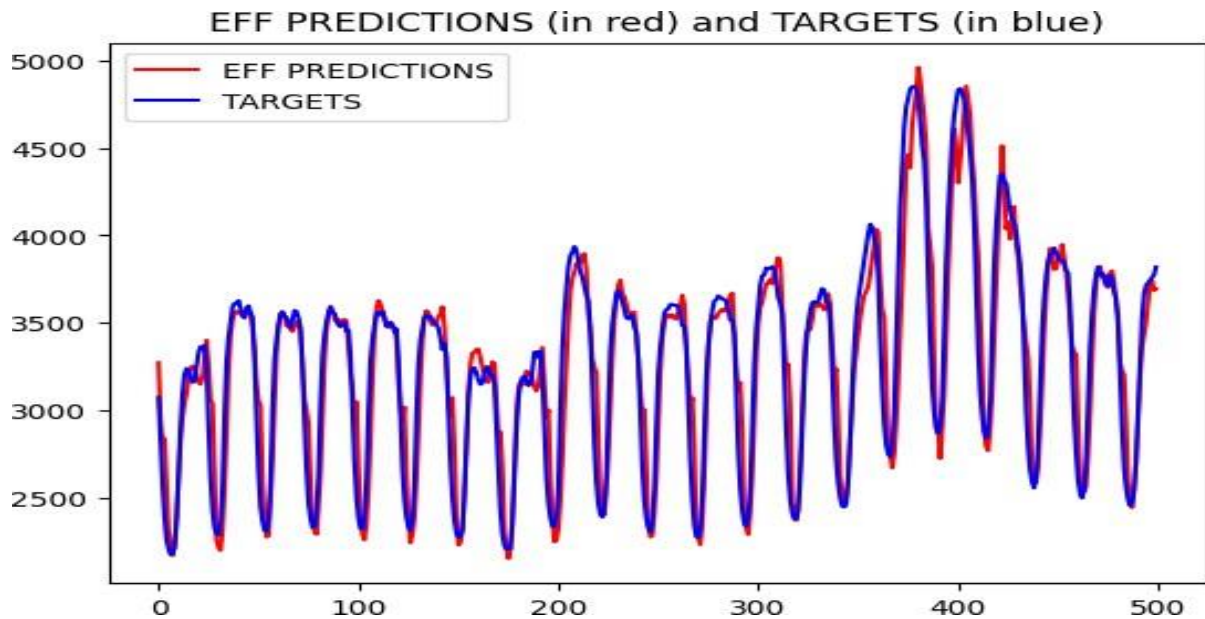
h2) Time series plot of all values predicted by the model on the test set:



i2) Time series plot of the corresponding targets:



j2) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL



k)

$$PMAE = \frac{(true)MAE \text{ obtained by the model on the test set}}{\text{Full range of elaad in the test set}}$$

$$PMAE = 172.75/3245$$

$$= 0.053$$

$$= 5.3\%$$

Part III-B: 2-Input Predictor (6 hour):

a) sequence_length = 24

b) COMPILE METHOD & FIT METHOD:

```
model.compile(loss='mae' , optimizer=opt , metrics=['mae'])

# Define callbacks using the best model only
best_model_callback = ModelCheckpoint("best_model.h5", monitor='val_loss',
    save_best_only=True)
early_stop = EarlyStopping(monitor='val_loss', patience=5)


# Train model
history = model.fit(train_dataset, epochs=60, batch_size = 128, validation
_data=val_dataset, callbacks=[best_model_callback])
```

c) model.summary()


Model: "sequential"

Layer (type)	Output Shape	Param #
=====	=====	=====
lstm (LSTM)	(None, 64)	17152
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 64)	4160
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
=====	=====	=====
Total params: 23,425		
Trainable params: 23,425		
Non-trainable params: 0		

d) TRAINING AND VALIDATION LOSS



e) TRAINING AND VALIDATION MAE



f) Final Training and Validation LOSS & MAE:

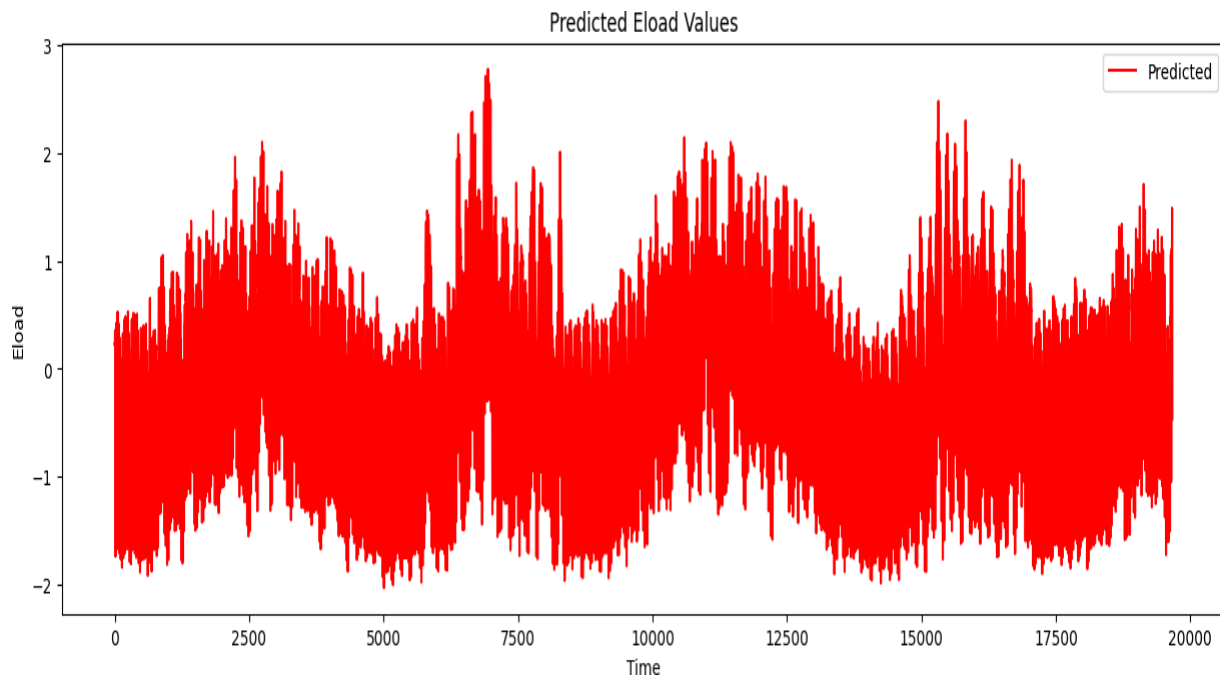
loss: 0.2723 - mae: 0.2723 - val_loss: 0.2335 - val_mae: 0.2335

Training Loss : 0.2723
Training MAE : 0.2723
Validation LOSS : 0.2335
Validation MAE : 0.2335


g) MAE obtained on the TEST SET:

Test MAE: 0.23702189326286316


h) Time series plot of all values predicted by the model on the test set:



i) Time series plot of the corresponding targets



j) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL




AFTER DOING DENORMALIZATION


g2) True MAE obtained on the TEST SET:

EFFECTIVE real-scale MAE: 158.03


h2) Time series plot of all values predicted by the model on the test set:



i2) Time series plot of the corresponding targets:



j2) Overlay plot of the predictions (red solid line) and the targets (blue solid line) for samples 6000 to 6500 OF THE PREDICTIONS FROM THE MODEL



k)

$$PMAE = \frac{(true)MAE \text{ obtained by the model on the test set}}{\text{Full range of elaad in the test set}}$$

$$PMAE = 158.03/3245$$

$$= 0.048$$

$$= 4.8\%$$

CONCLUSIONS

What I learned & difficulties faced:

Initially I was not able to get the better training and validation loss. After increasing the number of epochs and passing a learning rate argument to the optimizer, I was able to get the better results and obtained the best models after tuning a lot of hyperparameters. I was able to achieve finally Real-Scale of MAE of 158.03.

The difficulties I faced initially was without normalizing the targets I am not able to train and do predictions. The methodology which was provided helps me in doing normalization to get the better results. I think that the data with these large values we need to preprocess the data necessarily in order to scale down the values or else there will be a lot of misinterpretations and will lead to wrong results and predictions. So additional preprocessing of the data was a lot beneficial.

The PMAE of 1 vs 2-input predictors with a 3-hour horizon was 6.1 & 5.3, I observed some reduce in PMAE if increase the sequence length might get the better PMAE. I am still trying to reduce the PMAE of both the predictors. And the PMAE of 1 vs 2-input predictors with a 6-hour horizon was 5.9 & 4.8, I was able to get the best results for this predictors, So I can conclude that from my results I trained, the model was able to perform better when providing predictions of load 6 hours into the future.

